Android Malware Detection using Deep Learning on API Method Sequences

ElMouatez Billah Karbab, Mourad Debbabi, Abdelouahid Derhab, Djedjiga Mouheb

Concordia University, Concordia University, King Saud University, University of Sharjah

Abstract

Android OS experiences a blazing popularity since the last few years. This predominant platform has established itself not only in the mobile world but also in the Internet of Things (IoT) devices. This popularity, however, comes at the expense of security, as it has become a tempting target of malicious apps. Hence, there is an increasing need for sophisticated, automatic, and portable malware detection solutions. In this paper, we propose MalDozer, an automatic Android malware detection and family attribution framework that relies on sequences classification using deep learning techniques. Starting from the raw sequence of the app's API method calls, MalDozer automatically extracts and learns the malicious and the benign patterns from the actual samples to detect Android malware. MalDozer can serve as a ubiquitous malware detection system that is not only deployed on servers, but also on mobile and even IoT devices. We evaluate MalDozer on multiple Android malware datasets ranging from 1K to 33K malware apps, and 38K benign apps. The results show that MalDozer can correctly detect malware and attribute them to their actual families with an *F1-Score* of 96% − 99% and a *false positive* rate of 0.06% − 2%, under all tested datasets and settings.

Keywords: Mobile, Android, Malware, IoT, Deep Learning

1. Introduction

Mobile apps have become an inherent part of our everyday life since many of the services are provided to us through mobile apps. The latter change the way we communicate, as they are installed in most cases on smart devices. In contrast to personal computers, smart devices are equipped with sophisticated sensors, from cameras and microphones to gyroscopes and GPS [1]. These various sensors open a whole new world of applications for end-users [1], and generate huge amounts of data, which contain highly sensitive information. Consequently, this raises the need for security solutions to protect users from malicious apps, which exploit the sophistication of the smart devices and their sensitive data. On the other hand, the Internet of Things (IoT) smart systems have become equally, if not more, important than the mobile ones: (i) IoT systems are not only installed on conventional devices such as phones but are also considered in critical systems such as industrial IoT devices [2][3]. (ii) According to Ericsson [4], the number of IoT devices is expected to surpass the number of mobile devices by 2018 and could reach 16 billion by 2021. In this setting, security solutions should defend against malicious apps targeting both mobile and IoT devices. Android OS is phenomenally growing by powering a vast spectrum of smart devices. It has the biggest share in the mobile computing industry with 85% in 2017-Q1

[5] due to its open-source distribution and sophistication. Besides, it has become not only the dominant platform for mobile phones and tablets but is also gaining increasing attention and penetration in the IoT realm [6],[7],[8]. In this context, Google has launched Android Things [9], an Android OS for IoT devices, where developers benefit from the mature Android stack to develop IoT apps targeting thin devices [10], [9], [11], [12]. Therefore, protecting Android devices from malicious apps is of parmount importance.

Problem Statement. To address the above challenges, there is a clear need for a solution that defends against malicious apps in mobile and IoT devices with specific requirements to overcome the limitations of existing Android malware detection systems. First, the Android malware detection system should ensure a high accuracy with minimum false alarms. Second, it should be able to operate at different deployment scales: (i) Server machines, (ii) Personal machines, (iii) Smartphones and tablets, and (iv) IoT devices. Third, detecting that a given app is malicious may not be enough, as more information about the threat is needed to prioritize the mitigation actions. The type of attack could be crucial to prevent the intended damage. Therefore, it is essential to have a solution that goes a step further and attributes the malware to a specific family, which defines the potential threat that our system is exposed to. Finally, it is necessary to minimize manual human intervention to the largest extent and make the detection dependent mainly on the app sample for automatic feature extraction and pattern recognition. As malicious apps are quickly getting

stealthier, the security analyst should be able to catch up with this pace. This is due to the fact that for every new malware family, a manual analysis of the samples is required to identify its pattern and features that distinguish it from benign apps.

Solution. In this paper, we propose MalDozer, a simple, yet effective and efficient framework for Android malware detection based on sequences mining using neural networks. MalDozer framework is based on an artificial neural network that takes, as input, the raw sequences of API method calls, as they appears in the DEX file, to enable malware detection and family attribution. During the training, MalDozer can automatically recognize malicious patterns using only the sequences of raw method calls in the assembly code. MalDozer achieves a high accuracy in malware detection under multiple datasets, including Malgenome [13] (1K samples), Drebin [14] (5.5K samples), our MalDozer dataset (20K samples), and a merged dataset of 33K malware samples. Additionally, 38K benign apps downloaded from Google Play [15] are also used in the evaluation. MalDozer achieves an F1-score between 96% and 99% in the detection task. Furthermore, using the same datasets, MalDozer can correctly attribute the Android malware to the actual family with an F1-score between 96% and 98% in the family attribution task. MalDozer is both effective and also efficient. We evaluate the efficiency of MalDozer under multiple deployment architectures, ranging from high-end servers to very small IoT devices [12]. The results of our evaluation confirm that MalDozer can efficiently run on all these devices. The key idea of MalDozer relies on using neural networks on the API assembly method invocations to identify Android malware. More precisely, the input of MalDozer is the sequences of the API method calls as they appear in the *DEX* file, where a sequence represents the Android app. First, we map each method in the sequence invocation to a fixed length high-dimensional vector that semantically represents the method invocation [16] and replace the sequence of the Android app methods by a sequence of vectors. Afterward, we feed the sequence of vectors to a neural network with multiple layers. In this paper, we make the following contributions:

- MalDozer, a novel, effective, and efficient Android malware detection framework using the raw sequences of API method calls based on neural networks. We take a step beyond malware detection by attributing the detected Android malware to its family with a high accuracy.

- We propose an automatic feature extraction technique during the training using *method embedding*, where the input is the raw sequence of API method calls, extracted from DEX assembly.

- We conduct an extensive evaluation on different datasets real Android malware and benign apps. The

results demonstrate that MalDozer is very efficient and effective. It is also resilient against API evolution over time and against changing the order of API method calls. Additionally, MalDozer could be deployed and run properly, at various scales.

2. Background

In this section, we provide the necessary background that is relevant to our framework. We start by defining the cornerstone of MalDozer, namely neural network, and why it is interesting in the context of Android malware detection (Section 2.1). Afterward, we present the threat model as well as the assumptions considered in MalDozer design (Section 2.2). Next, we enumerate the main use cases of MalDozer framework (Section 2.3).

2.1. Deep Learning and Neural Network

A neural network is a machine learning computation model, which relies on a large number of neural units. The latter are approximate abstractions of the brain neurons, which could solve a very complex problem using highly dense neurons connected to each other by axons. Typically, Artificial Neuron Network (ANN) is composed of multiple layers, where each layer has many artificial neurons. The first layer is the input layer, and the last layer is the output one. The rest of the layers are called hidden layers. Notice that the neurons in each layer i are connected to layer $i+1$, but the connection method could differ from a model to another. To this end, in the deep learning terminology, a neural network consists of multiple hidden layers, i.e., the more layers there are, the deeper the neural network is. The conventional machine learning methods are limited by the manually-crafted features from the raw data. Here, the security expert analyzes the malicious apps and extracts the relevant features. The latter will be fed to a classifier to produce a learning model. The main advantage of a neural network is that it could automatically learn the representation (features) from the raw data to perform the detection task. In this paper, we aim at taking a step further towards Android malware detection with automatic representation learning. To achieve this aim, we leverage deep learning techniques and only consider the raw API method calls from Android DEX files for the purpose of malware detection and attribution with automatic feature extraction.

2.2. Threat Model and Assumptions

We position MalDozer as an anti-malware system that detects Android malware and attributes it to a known family with a high accuracy and minimal false positive and negative rates. We assume that the analyzed Android apps, whether malicious or benign, are developed mainly in Java or any other language that is translated to DEX bytecode. Therefore, Android apps developed by other

means, e.g., web-based, are out of the scope of the current design of MalDozer. Also, we assume that apps' core functionalities are in the DEX bytecode and not in C/C++ native code [17], i.e., the attacker is mainly using the DEX bytecode for the malicious payload. Furthermore, we assume that MalDozer detection results could not be affected by malicious activities. In the case of a server, Android malicious apps have no effect on the server system. However, in the case of deployment on infected mobiles or IoT devices, MalDozer should be protected from malicious activities to avoid tampering its results.

2.3. Usage Scenarios

The effectiveness of MalDozer, i.e., its high accuracy, makes it a suitable choice for malware detection in large-scale app store systems, especially that its update only requires very minimal manual intervention. We only need to train MalDozer model on new samples without a *feature engineering*, since MalDozer can automatically extract and learn the malicious and benign features during the training. Notice that MalDozer could detect unknown malware based on our evaluation as presented in Section 7. Furthermore, due to the efficiency of MalDozer, it could be deployed on mobile devices such as phones and tablets. As for mobile devices, MalDozer acts as a detection component in the anti-malware system, where the goal is to scan new apps. The family attribution is very handy when detecting new malware apps. Indeed, MalDozer helps the anti-malware system to take the necessary precautions and actions based on the malware family, which could have some specific malicious threats such as ransomware. It is also important to mention that we were able to run MalDozer on resource-limited IoT devices considered by Android Things such as Raspberry PI [12].

2.4. Android Architecture

Android has been settled by the Android Open Source Project (AOSP) team, maintained by Google and supported by the Open Handset Alliance (OHA) [18]. It encompasses the Original Equipment Manufacturers (OEMs), chip-makers, carriers and application developers. Android apps are written in Java. However, the native code and shared libraries are generally developed in C/C++ [17]. The current Android architecture [19] consists of a Linux kernel, which is designed for an embedded environment consisting of limited resources. On top of the Linux kernel, there is a Hardware Abstraction Layer (HAL), which provides standard interfaces that expose device hardware capabilities to the higher-level Java API framework, by allowing programmers to create software hooks between the Android platform stack and the hardware. There is also Android Runtime (ART), which is an application runtime environment used by the Android OS. It replaced Dalvik starting from Android 5.0. ART translates the app's bytecode into native instructions that are later executed by the

device's runtime environment. ART introduces the ahead-of-time (AOT) compilation feature, which allows compiling entire applications into native machine code upon their installation. The native libraries developed in C/C++ support high-performance third-party reusable shared libraries. The Java API Framework provides APIs form the building blocks the user need to create Android apps. The System Apps, are the apps hat are included within the system, and are required to make the device run.

2.4.1. Android APK Format

Android Application Package (APK) is the file format adopted by Android for apps distribution and installation. It comes as a *ZIP* archive file, which contains all the components needed to run the app. By analogy, APK files are similar to Windows *EXE* installation files or Linux RPM/DEB files. The APK package is organized into different directories (namely **lib**, **res**, and **assets**) and files (namely **AndroidManifest.xml** and classes.dex). More precisely, i) The **AndroidManifest.xml** file contains the app meta-data, e.g., name, version, required permissions, and used libraries. ii) The classes.dex file contains the compiled Java classes. iii) The **lib** directory stores C/C++ native libraries [17]. iv) The resources directory (**res**) contains the non-source code files, such as video, image, and audio files, which are packaged during compilation.

3. Methodology

In this section, we present MalDozer framework and its components (Figure 1). MalDozer has a simple design, where a minimalistic preprocessing is employed to get the assembly methods. As for the feature extraction (representation learning) and detection/attribution, they are based on the actual neural network. This permits MalDozer to be very efficient with fast preprocessing and neural network execution. Since MalDozer is based on a supervised machine learning, we first need to train our model. Afterward, we deploy this model along with a preprocessing procedure on the targeted devices. Notice that the preprocessing procedure is common between the training and the deployment phases to ensure the correctness of the detection results (Figure 1).

1- Extraction of API Method Calls. MalDozer workflow starts by extracting the sequences of API calls from Android app packages, in which we consider only the *DEX* file. We disassemble the classes.dex to produce the Dalvik VM assembly. Our goal is to formalize the assembly to keep the maximum raw information with minimum noise. Notice here that we could use Android APIs (such as `android/net/ConnectivityManager` in Figure 2) instead of permission to have a granular view that helps distinguishing a malware app.

However, quantifying Android API could be noisy because there are plenty of common API calls shared between apps. Some solutions tend to filter only dangerous
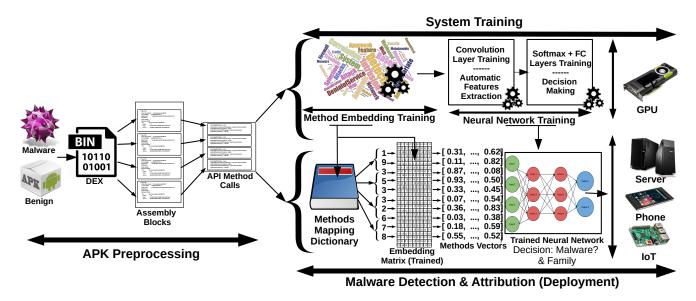
Figure 1: Approach Overview

```
android/net/ConnectivityManager
android/net/ConnectivityManager
android/telephony/SmsManager
android/telephony/SmsManager
android/location/LocationManager
android/location/LocationManager
```

Figure 2: Android API from a Malware Sample

APIs and use them for detection. In this case, we require a manual categorization of dangerous APIs. Moreover, Android API gives an abstract view of the actual malicious activity that could deceive the malware detection. For this reason, we leverage Android API method calls as `android/net/ConnectivityManager;-> getNetworkInfo` in Figure 3. By doing so, the malware detector will have a more granular view of the app activity. In our case, we address this problem from another angle; we treat Android apps as a sequence of API method calls. We consider all the API calls with no filtering, where the order is part of the information we use to identify malware. It represents the temporal relationship between two API method calls (in a given basic block), and defines the intended sub-tasks of the app. The sequence of API method calls preserves the temporal relationship over individual basic blocks of the linear disassembly and ignores the order between these blocks. The obtained result is a merged sequence (Figure 1).

```
android/net/ConnectivityManager;->getNetworkInfo
android/net/ConnectivityManager;->getAllNetworkInfo
android/telephony/SmsManager;->sendTextMessage
android/telephony/SmsManager;->sendMultipartTextMessage
android/location/LocationManager;->getLastKnownLocation
android/location/LocationManager;->getBestProvider
```

Figure 3: Granular View using API Method Calls

In other words, a *DEX* file, denoted by cd, is composed of a set of n compiled Java classes, $cd = \{cl_1, \cdots, cl_n\}$. Each Java class cl_i is, in turn, composed of a set of m methods, which are basic blocks, $cl_i = \{mt_1^i \cdots, mt_m^i\}$. By going down to the API method level, mt_j^i is a sequence of k API method calls, Formally $mt_j^i = (P_1^{i,j}, \cdots, P_k^{i,j})$, where $P_l^{i,j}$ is the l^{th} API method call in method mt_j^i.

Algorithm 1: Extraction

Input : cd: Java Assembly
Output: $MSeq$: Methods Sequence

begin
    $MSeq = \texttt{EmptyList}();$

    foreach $cl \in cd$ **do**
        foreach $mt \in cl$ **do**
            foreach $P \in mt$ **do**
                $MSeq.Add(p);$
            end
        end
    end

    return $MSeq$;
end

2- Discretization of API Method Calls. In this step, we discretize the sequences of API method calls that are in an Android app (Algorithm 2). More precisely, we replace each API method with an identifier, resulting in a sequence of numbers. We also build a dictionary that maps each API call to its identifier. Notice that in the current implementation, the mapping dictionary is deployed with the learning model to map the API calls of the analyzed apps. In the deployment, we could find unknown API calls

related to third party libraries. To overcome this problem: (i) We consider a big dataset that covers most of the API calls. (ii) In the deployment phase, we replace unknown API calls with fixed identifiers. Afterward, we unify the length of the sequences L (hyperparameter) and pad a given sequence with zeros if its length $l < L$.

Algorithm 2: Discretization

Input : $MSeq$: Methods Sequence
            $MapDict$: Mapping Dict
Output: $DSeq$: Discrete Sequence
begin
    $DSeq = \texttt{EmptyList}()$;
    foreach $P \in MSeq$ **do**
        if $m \in MapDict.Keys()$ **then**
            $Dvalue \leftarrow MapDict[P]$;
            $DSeq.Add(Dvalue)$;
        else
            $Dvalue \leftarrow 0$;
            $DSeq.Add(Dvalue)$;
        end
    end
    return $DSeq$;
end

3- Unification of the Sequences' Size. The length of the sequences varies from one app to another. Hence, it is important to unify the length of the sequences. There are two cases depending on the length of the sequence and the hyper-parameter. We choose a uniform sequence size as follows: i) If the length of a given sequence is greater than the uniform sequence size L, we take only the first L items to represent the apps. ii) In case the length of the sequence is less than L, we pad the sequence with zeros. It is important to mention that the uniform sequence size hyper-parameter has an influence on the accuracy of MalDozer. A simple rule is that the larger is the size, the better is, but this will require a lot of computation power and a long time to train the neural network.

3- Generation of the Semantic Vectors.. The identifier in the sequences needs to be shaped to fit as input to our neural network. This could be solved by representing each identifier by a vector. The question that arises is *how are such vectors produced?* A straightforward solution is to use one-hot vectors, where a vector has one in the interface value row, and zero in the rest. Such a vector is very sparse because its size is equal to the number of API calls, which makes it impractically and computationally prohibitive for the training and the deployment. To address this issue, we resort to a dense vector that uses a continuous space. These vectors are semantically related, and we could express their relation by computing a distance. The smaller the distance is, the more related the vectors are (i.e., the API calls). We describe word embedding in Section 4. The output of this step is sequences of vectors

for each app that keeps the order of the original API calls; each vector has a fixed size K (hyper-parameter).

4- Prediction using a Neural Network. The final component in MalDozer framework is the neural network, which is composed of several layers. The number of layers and the complexity of the model are hyper-parameters. However, we aim to keep the neural network model as simple as possible to gain in the execution time during its deployment, especially on IoT devices. In our design, we rely on the convolution layers [20] to automatically discover the pattern in the raw method calls. The input to the neural network is a sequence of vectors, i.e., a matrix of $L \times K$ shape. In the training phase, we train the neural network parameters (layers weight) based on the app vector sequence and its labels: (i) malware or benign for the detection task, and (ii) malware families for the attribution task. In the deployment phase, we extract the sequence of methods and use the embedding model to produce the vector sequence. Finally, the neural network takes the vector sequence to decide about the given Android app.

4. MalDozer Method Embedding

The neural network takes vectors as input. Therefore, we represent our Android API method calls as vectors. As a result, we formalize an Android app as a sequence of vectors with fixed size (L). We could use one-hot vector. However, its size is the number of unique API method calls in our dataset. This makes such a solution not scalable to large-scale training. Also, the word embedding technique outperforms the results of the one-hot vector technique in our case [16], [21], [20]. Therefore, we seek a compact vector, which also has a semantic value. To fulfill these requirements, we choose the word embedding techniques, namely, word2vec [16] and GloVe [21]. Our primary goal is to have a dense vector for each Android API method that keeps track of its contexts in a large dataset of Android apps. Thus, in contrast with one-hot vectors, each word embedding vector contains a numerical summary of the Android API call meaning representation. Moreover, we could apply geometric techniques on the API call vectors to measure the semantic relationship between their functionalities, i.e., developers tend to use certain API method calls in the same context. In our context, we learn these vectors from our dataset that contains benign and malicious apps by using word2vec [16]. The latter is a computationally efficient predictive model from learning word embedding vectors, which are applied on the raw Android API method calls. The output obtained from training the embedding word model is a matrix $K \times A$, where K is the size of the embedding vector, and A is the number of unique Android API method calls. Both K and A are hyper-parameters; we use $K = 64$ in all our models. In contrast, the hyper-parameter A is a major factor in the accuracy of MalDozer. The more API calls we consider, the more accurate and robust our model is. Notice that,

our word embedding is trained along with the neural network, where we tune both of them for a given task such as detection. Despite that, it can be trained separately to generate the embedding word vector independently of the detection task. In the deployment phase (Figure 1), MalDozer uses the word embedding model and looks up for each API method call identifier to find the corresponding embedding vector .

5. MalDozer Neural Network

MalDozer neural network is inspired by [20], where the authors use a neural network for sentence classification task such as sentiment analysis. The proposed architecture shows high results and outperforms many of state-of-the-art benchmarks with a relatively simple neural network design. Here, we raise the following questions: *Why could such a Natural Language Processing (NLP) model be useful in Android malware detection?* And *why do we choose to build it on top of this design [20]?* We formulate our answers as follows: i) NLP is a challenging field where we deal with text. So, there is an enormous number of vocabularies; also we could express the same meaning in different ways. We also have the same semantics with many combinations of words, which we call the *natural language obfuscation*. In our context, we deal with sequences of Android API method calls and want to find the combination of patterns of method calls, which produces the same (malicious) activity. We use the API method calls as they appear in the binary, i.e., there is a temporal relationship between API methods in basic blocks but we ignore the order among these blocks. By analogy to NLP, the basic blocks are the sentences and the API method calls are the words. Further, the app (paragraph) is a list of basic blocks (unordered sentences). This task looks easier compared to the NLP one because of the huge difference in the vocabulary, i.e., the number of Android API method calls is significantly less than the number of words in natural language. Also, the combination in the NLP is much complex compared to Android API calls. ii) We choose to use this model due to its efficiency and ability to run our model on resource-constrained devices. Table 1 depicts the neural network architecture of MalDozer's detection and attribution tasks. Both networks are very similar; the only notable difference is in the output layer. In the detection task, we need only one neuron in the output layer because the network decides whether the app is malware or not. As for the attribution task, there are multiple neurons, one for each Android malware family. Having the same architecture for the detection and attribution makes the development and the evaluation of a given design more simple. Because the network architecture achieves good results in one task, it will have very similar results in the other one. As presented in Figure 4, the first layer is a convolution layer [20] with rectified linear unit (ReLU) activation function ($f(x) = max(0, x)$). Afterward, we use global max pool [20] and connect it to a

fully-connected layer. Notice that in addition to Dropout [22] used to prevent overfitting, we also utilize Benchnormalization [22] to improve our results. Finally, we have an output layer, where the number of neurons depends on the detection or attribution tasks.

#	Layers	Options	Activ
1	Convolution	Filter=512, FilterSize=3	ReLU
2	MaxPooling	/	/
3	FC	#Neurons=256, Dropout=0.5	ReLU
4	FC	#Neurons={1,#Families[1]}	Softmax

[1] The number of malware families in the training dataset.

Table 1: MalDozer Malware Neural Network

	Server (1/2)	Laptop	Raspberry PI 2
GPU	TITAN X / no	no	no
CPU	Intel E5-2630	Intel T6400	ARM Cortex A7
RAM	128GB	3GB	1GB

Table 2: Hardware Specifications

6. Implementation

In this section, we present the software & hardware components of MalDozer evaluation.

Software. We implement MalDozer using *Python* and *Bash* scripting languages. First, Python zip library extracts the *DEX* file from the *APK* file. We use *dexdump* command-line to produce the assembly from the DEX file. *Dexdump* is available through the Android SDK, but in the case of Raspberry PI, we built it from its source code. Regular expressions are employed to extract API method calls from the assembly. To develop the neural network, we use Tensorflow [23]. Notice that there is no optimization in the preprocessing; in the run-time evaluation, we use only a single thread app.

Hardware. To evaluate the efficiency of MalDozer, we evaluate multiple types of hardware, as shown in Table 2, starting from servers to *Raspberry PI* [12]. For training, the Graphic Processing Unit (GPU) is a vital component because the neural network training needs immense computational power. The training takes hours under *NVIDIA TitanX*. However, the deployment could be virtually on any device including IoT devices. To this end, we consider *Raspberry PI* as IoT device because it is one of the hardware platforms supported by Android Things [9]. We also use low-end laptops in our evaluation, as shown in Table 2.
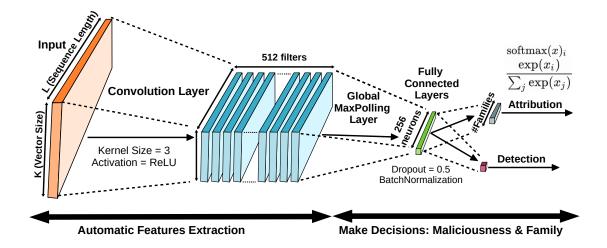
Figure 4: Neural Network Architecture

7. Evaluation

In this section, we conduct our' evaluation using different datasets that primarily cover the following performance aspects:

- (I) *Detection Performance*: We evaluate how effectively MalDozer can distinguish between malicious and benign apps in terms of F1-measure, precision, recall, and false positive rate.

- (II) *Attribution Performance*: We evaluate how effectively MalDozer can correctly attribute a given malicious app to its malware family.

- (III) *Runtime Performance*: We measure the preprocessing and the detection runtime on different types of hardware.

7.1. Evaluation Metrics

The evaluation results are presented under the following metrics:

- *True positives* (TP): This metric measures the number of malicious apps that are successfully detected.

- *False negatives* (FN): This metric measures the number of malicious apps that are incorrectly classified.

- *False positives* (FP): This metric measures the number of benign apps that are incorrectly classified.

- *Precision* (P): It is the percentage of positive prediction, i.e., the percentage of the detected malware out of all sample apps. Formally, $P = \frac{TP}{TP+FP}$

- *Recall* (R): It is the percentage of correct malicious apps detected out of all malware samples. Formally, $R = \frac{TP}{TP+FN}$

- *F1-Score* (F1): It is a measure that considers precision and recall. Formally, $F1 = 2 \times \frac{P \times R}{P+R}$.

We also measure *False Positive Rate* (FPR), *False Negative Rate* (FNR) and *Accuracy* (ACC), which are given as follows:

$$FPR = \frac{FP}{FP+TP} \;,\; FNR = \frac{FN}{FN+TP} \;,\; ACC = \frac{TP+TN}{P+N}$$

7.2. Datasets

In our evaluation, we have two main tasks: i) Detection, which aims at checking if a given app is malware or not, ii) Attribution, which aims at determining the family of the detected malware. We conduct the evaluation experiments under two types of datasets: i) Mixed dataset, which contains malicious apps and benign apps, as presented in Table 3. ii) Malware dataset, which contains only malware, as shown in Table 4. As for the malware dataset, we leverage reference datasets such as *Malgenome* [13] and *Drebin* [24]. We also collect two other datasets from different sources, e.g., *virusshare.com*, *Contagio Minidump* [25]. The total number of malware samples is $33K$, including Malgenome and Drebin datasets. As for the attribution task, we use only malware from the previous datasets, where each family has at least 40 samples, as presented in Tables 13, 14, 15. To this end, we propose MalDozer dataset, as in Table 13, which contains 20K malware samples from 32 malware families. We envision to make MalDozer dataset available upon request for the research community. The benign app samples have been collected from *Playdrone* dataset [26]. We leverage the top $38K$ apps that are ranked by the number of downloads.

7.3. Malware Detection Performance

We evaluate MalDozer on different cross-validation settings, two, three, five and ten-fold, to examine the detection performance under different training/test set percentages ($50\%, 66\%, 80\%, 90\%$) from the actual dataset (10 training epochs). Table 5 depicts the detection results

Dataset	#Malware	#Benign	Total
Malgenome	1,258	37,627	38.885
Drebin	5,555	37,627	43,182
MalDozer	20,089	37,627	57,716
All	33,066	37,627	70,693

Table 3: Datasets for Detection Task

Dataset	#Malware	#Family
Malgenome	985	9
Drebin	4,661	20
MalDozer	20,089	32

Table 4: Datasets for Attribution Task

on Malgenome dataset. MalDozer achieves excellent results, F1-Score=99.84%, with a small *False Positive Rate* (FPR), 0.04%, despite the unbalanced dataset, where benign app samples are the most dominant in the dataset. The detection results are similar under all cross-validation settings. Table 6 presents the detection results on Drebin dataset, which are very similar to the Malgenome ones. MalDozer reaches F1-Score=99.21%, with FPR=0.45%. Similar detection results are shown in Table 7 on MalDozer dataset (F1-Score=98.18% and FPR=1.15%). Table 8 shows the results related to all datasets, where MalDozer achieves a good result (F1-Score=96.33%). However, it has a higher false positive rate compared to the previous results (FPR=3.19%). This leads us to manually investigate the false postives. We discover, by correlating with *virusTotal.com*, that several false positive apps are already detected by many vendors as malware.

	F1%	P%	R%	FPR%
2-Fold	99.6600	99.6620	99.6656	0.06
3-Fold	98.1926	98.6673	97.9812	1.97
5-Fold	99.8044	99.8042	99.8045	0.09
10-Fold	**99.8482**	**99.8474**	**99.8482**	**0.04**

Table 5: Detection on Malgenome Dataset

	F1%	P%	R%	FPR%
2-Fold	98.8834	98.9015	98.9000	0.13
3-Fold	99.0142	99.0130	99.01579	0.51
5-Fold	99.1174	99.1173	99.1223	0.31
10-Fold	**99.2173**	**99.2173**	**99.2172**	**0.45**

Table 6: Detection on Drebin Dataset

7.3.1. Unknown Malware Detection

Although MalDozer demonstrates very good detection results, some questions still arise: (i) *Can MalDozer detect samples of unknown malware families?* And (ii) *How*

	F1%	P%	R%	FPR%
2-Fold	96.8576	96.9079	96.8778	1.01
3-Fold	97.6229	97.6260	97.6211	2.00
5-Fold	97.7804	97.7964	97.7753	2.25
10-Fold	**98.1875**	**98.1876**	**98.1894**	**1.15**

Table 7: Detection on MalDozer Dataset

	F1%	P%	R%	FPR%
2-Fold	96.0708	96.0962	96.0745	2.53
3-Fold	95.0252	95.0252	95.0278	4.01
5-Fold	96.3326	96.3434	96.3348	2.67
10-Fold	**96.2958**	**96.2969**	**96.2966**	**3.19**

Table 8: Detection on All Dataset

many samples are needed for a given family to achieve a good accuracy? To answer these questions, we conduct the following experiment on Drebin mixed dataset (Malware + Benign), where we focus on top malware families (i.e., BaseBridge, DroidKungFu, FakeInstaller, GinMaster, Opfake, Plankton). For each family, we train (5 epochs) our model on a subset dataset, which does not include samples of that family. These samples are used as a test set. Afterward, we train with few samples from the family and evaluate the model on the rest of them. Progressively, we add more samples to the training and assess the accuracy of our model on detecting the rest of the family samples. Answering the above questions: *(i) Can MalDozer detect unknown malware family samples?* Yes, Figure 5 shows the accuracy versus the number of samples in the training dataset. We see that MalDozer (zero sample vs. accuracy) could detect the unknown malware family sample without previous training. The accuracy varies from 60% to 90%. *(ii) How many samples for a given family to achieve a good accuracy?* MalDozer needs only about 10 to 20 samples to reach 90% (Figure 5). In the case of DroidKungFu, MalDozer needs 20 samples to reach 90%. Considering 10 to 20 samples from a malware family is relatively a small number to get high results.

7.3.2. Resiliency Against API Evolution over Time

As we have seen in the previous section, MalDozer could detect new malware samples from unknown families using samples from *Drebin* dataset collected in the period of 2011/2012. In this section, we aim to answer the following quastion: *Can MalDozer detect malicious and bengin apps collected in different years?* To answer this question, we evalute MalDozer on four datasets collected from [27] of four consecutive years: 2013, 2014, 2015, and 2016, as shown in Table 9, where, we train MalDozer in one year dataset and test it on the rest of the datasets. The results show that MalDozer detection is more resilient to API evolution over time compare to [28], as presented in Figure 6. Starting with 2013 dataset (Figure 6(a)), we train MalDozer on 2013 samples and evaluate it on 2014, 2015,

(a) BaseBridge (b) DrKungFu (c) FakeInst (d) GinMaster
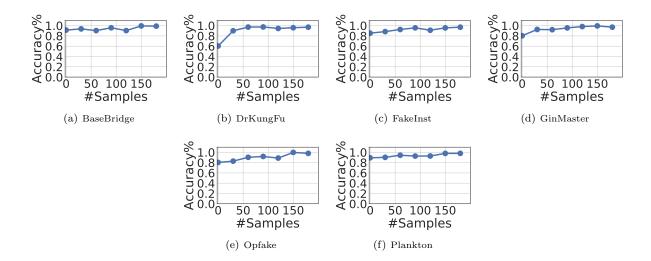
(e) Opfake (f) Plankton

Figure 5: Evaluation of Unknown Malware Detection

and 2016 ones. We notice a high detection rate in 2014 dataset since it is collected in the consecutive year of the training datset. However, the detection rate decreases in 2015 and 2016 datasets but it is above an acceptable detection rate (F1-Score=70%). Similarly, we obtained the results of 2014 dataset, as depicted in Figure 6(b). Also, training MalDozer on 2015 or 2016 datasets exhibits very good results under all the datasets collected in other years, where we reach F1-Score=90-92.5%.

Dataset	2013	2014	2015	2016
#Malware	10k	10k	10k	10k
#Bengin	10k	10k	10k	10k

Table 9: Datasets for API Evolution Over Time Resiliency

7.3.3. Resiliency against changing the order of API methods

In this section, we evaluate the robustness of MalDozer against changing in the order of API method calls. The latter could change for various reasons, for example: (i) We could use different dissassembly tools in the production, (ii) A malware developer could repackage the same malicious app multiple times. The previous scenarios could lead to losing the temporal relations among the API calls. In case of the malware developer, he/she will be limited by keeping the same malicious semantics in the app. To validate the robustness of MalDozer against such methods order, we conduct the following experiment. First, we train our model on the training dataset. Afterward, we randomly shuffle the sequence of API method calls in the test dataset. We divide the testing app sequence into N blocks, then shuffle them and evaluate the F1-Score. We repeat until N is equal to the number of sequences, i.e., one API call in each block. The result of this experiment is shown in Figure 7. The latter depicts the F1-Score versus

the number of blocks, starting with four blocks and ending with 15K blocks, where each block contains one API call. Figure 7 demonstrates the resiliency of MalDozer against changing the order of API method calls. We observe that even with completely random individual API method calls, MalDozer achieves 93%.

7.4. Family Attribution Performance

Family attribution is an important task for Android security, where MalDozer distinguishes itself from the existing malware detection solutions, since only few solutions provide this functionality. Starting with Malgenome dataset, MalDozer achieves a very good result, i.e., F1-Score of 99.18%. Similarly, MalDozer reaches an F1-Score of 98% on Drebin dataset. The results per malware family attribution performance for Malgnome and Drebin are presented in Tables 14 and 15. MalDozer achieves good results in the case of MalDozer dataset, F1-Score of 85%. Our interpretation of this result comes from Tables 13, 14 and 15, which depict the detailed results per malware family. For example, the family agent unveils poor results because of the misslabeling, since agent is a common name for many Android malware families. We believe that there is a lot of noise in the family labeling of the MalDozer dataset since we leverage only one security vendor for labeling. Despite this fact, MalDozer demonstrates acceptable results and robustness (Appendix Figure A.13).

	F1%	P%	R%
2-Fold	98.9834	99.0009	98.9847
3-Fold	98.9910	99.0026	98.9847
5-Fold	99.0907	99.1032	99.0862
10-Fold	**99.1873**	**99.1873**	**99.1878**

Table 10: Attribution on Malgenome

(a) 2013 Dataset

(b) 2014 Dataset

(c) 2015 Dataset

(d) 2016 Dataset

Figure 6: Detection vs Time

Figure 7: Shuffle Rate vs F1-Score

	F1%	P%	R%
2-Fold	98.1192	98.1401	98.1334
3-Fold	98.6882	98.6998	98.6912
5-Fold	98.5824	98.5961	98.5839
10-Fold	**98.5198**	**98.5295**	**98.5196**

Table 11: Attribution on Drebin

	F1%	P%	R%
2-Fold	89.3331	89.5044	89.3424
3-Fold	81.8742	82.7565	81.8109
5-Fold	83.8518	84.1360	84.0061
10-Fold	**85.5233**	**85.6184**	**85.8479**

Table 12: Attribution on MalDozer

7.5. Run-Time Performance

In this section, we evaluate the efficiency of MalDozer, i.e., the runtime during the deployment phase. We divide the runtime into two parts: i) *Preprocessing time*: the required time to extract and preprocess the sequences of Android API method calls. ii) *Detection time:* time needed to make the prediction about a given sequence of API

method calls. We analyze the detection time on the model complexity of different hardware. Figure 11(a) depicts the average preprocessing time along with its standard deviation, related to each hardware. The server machines and the laptop spend, on average, 1 second in the preprocessing time, which is very acceptable for production. Also, as mentioned previously, we do not optimize the current preprocessing workflow. In the IoT device [12], the preprocessing takes, on average, about 4 seconds, which is more than acceptable for such a small device. Figure 11(b) presents the detection time on average that is related to each hardware. First, it is noticeable that the standard deviation is very negligible, i.e., the detection time is constant for all apps. Also, the detection time is very low for all the devices. As for the IoT device, the detection time is only 1.3 seconds. Therefore, the average time that MalDozer needs to decide for a given app is 5.3 seconds on average in case of IoT device, as we know that the preprocessing takes most of the time (4/5.3). Here, we ask the following two questions: (i) *Which part in the pre-*

		Malware Family	#Sample	F1-Score
01		FakeInst	4822	96.15%
02		Dowgin	2248	84.24%
03		SmsPay	1544	81.61%
04		Adwo	1495	87.79%
05		SMSSend	1088	81.48%
06		Wapsx	833	78.85%
07		Plankton	817	94.18%
08		Agent	778	51.45%
09		SMSReg	687	80.61%
10		GingerMaster	533	76.39%
11		Kuguo	448	78.28%
12		HiddenAds	426	84.20%
13		Utchi	397	93.99%
14		Youmi	355	72.39%
15		Iop	344	93.09%
16		BaseBridge	341	90.50%
17		DroidKungFu	314	85.85%
18		SmsSpy	279	85.05%
19		FakeApp	278	93.99%
20		InfoStealer	253	82.82%
21		Kmin	222	91.03%
22		HiddenApp	214	76.71%
23		AppQuanta	202	99.26%
24		Dropper	195	77.11%
25		MobilePay	144	78.74%
26		FakeDoc	140	96.38%
27		Mseg	138	55.38%
28		SMSKey	130	81.03%
29		RATC	111	84.81%
30		Geinimi	106	95.58%
31		DDLight	104	90.55%
32		GingerBreak	103	84.87%

Table 13: MalDozer Android Malware Dataset

		Malware Family	#Sample	F1-Score
01		DroidKungFu3	309	99.83%
02		AnserverBot	187	99.19%
03		BaseBridge	121	98.37%
04		DroidKungFu4	96	99.88%
05		Geinimi	69	97.81%
06		Pjapps	58	95.65%
07		KMin	52	99.99%
08		GoldDream	47	99.96%
09		DroidDreamLight	46	99.99%

Table 14: Malgenome Attribution Dataset

processing needs optimization? (ii) *Does the preprocessing time depend on the size of APK or DEX file?* To answer these questions, we randomly select 1K benign apps and 1K malware apps. We measure the preprocessing time and correlate it with the size of APK and DEX files. Figure 8 shows the experimentation results in the case of the IoT device [12]. The scattered charts depict the preprocess-

		Malware Family	#Sample	F1-Score
01		FakeInstaller	925	99.51%
02		DroidKungFu	666	98.79%
03		Plankton	625	99.11%
04		Opfake	613	99.34%
05		GinMaster	339	97.92%
06		BaseBridge	329	97.56%
07		Iconosys	152	99.02%
08		Kmin	147	99.31%
09		FakeDoc	132	99.24%
10		Geinimi	92	97.26%
11		Adrd	91	96.13%
12		DroidDream	81	98.13%
13		Glodream	69	90.14%
14		MobileTx	69	91.97%
15		ExploitLinuxLotoor	69	99.97%
16		FakeRun	61	95.16%
17		SendPay	59	99.14%
18		Gappusin	58	97.43%
19		Imlog	43	98.85%
20		SMSreg	41	92.30%

Table 15: Drebin Attribution Dataset

ing time along with the size of the APK or DEX file for the mixed, only-benign, and only-malware datasets. From Figure 8, it is clear that the preprocessing time is linearly related to the size of the DEX file. We perform the same experiment on Server and Laptop, and we get very similar results, as shown in Figures 9 and 10. Finally, we notice that the size of benign apps tend to be bigger than the malicious apps. Thus, the preprocessing time of the benign apps is longer.

	#Params	F1%	Word2Vec Size
Model 01	6.6 Million	98.95	100k
Model 02	4.6 Million	95.84	70k
Model 03	3.4 Million	93.81	50k
Model 04	1.5 Million	90.08	20k

Table 16: Model Complexity vs Detection Performance

7.5.1. Model Complexity Evaluation

In this section, we examine the effect of model complexity on the detection time. By model complexity, we mean the number of parameters in the model, as depicted in Table 16. Many hyper-parameters could influence the complex nature of the model, but we primarily consider the word2vec embedding size. The latter is very important for the detection of the model, especially if we have a big dataset. Table 16 demonstrates the complexity of the model versus the F1-Score. It is noticeable that the larger the number of parameters is, the more its performance increases. Based on our observation, bigger models are more accurate and more robust to changes, as will be discussed
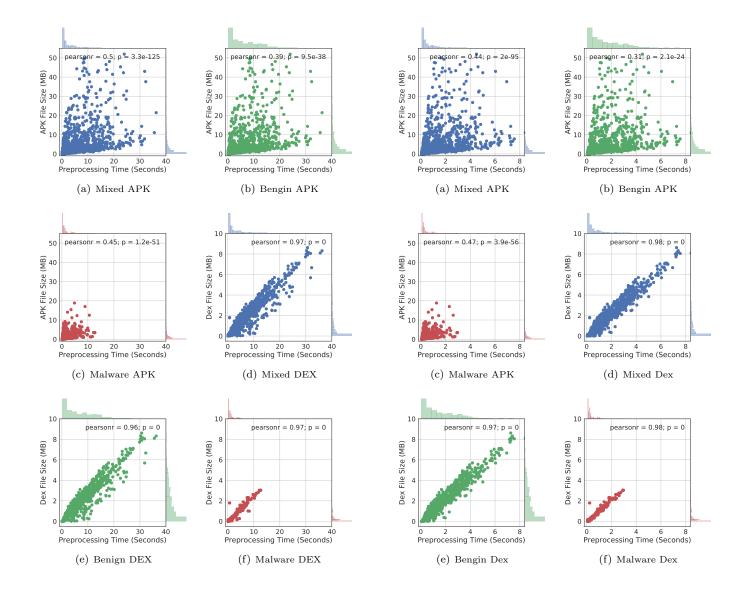
11

Figure 8: Preprocessing Time vs APK and DEX Sizes (IoT device)

Figure 9: Preprocessing Time vs APK and Dex Sizes (Laptop)

in Section 8. Finally, Figure 12 displays the execution time of the models in Table 16 on the IoT device. The detailed execution related to all the hardware is presented in Figure 12.

8. Discussion and Limitations

In this paper, we have explored a new approach to capture Android apps behaviors using neural networks on API method calls. This approach achieves highly accurate malware detection and family attribution. Our detection technique is sample-based, i.e., the system could automatically recognize patterns in the training phase of new malware as well as benign apps from raw sequences of API method calls. Therefore, this allows our system to catch up with the rapid evolution of Android OS and malicious techniques by training it on the raw sequence of API methods of new apps, which contain a lot of information about

the app's behaviors. Yet, this sequence is less affected by the obfuscation techniques. Furthermore, our work pushes toward portable detection solutions, i.e., the solution should be used in app stores, mobile or IoT devices. A portable solution is a step towards ubiquitous security that enhances small devices security. In this context, MalDozer could resist to certain obfuscation techniques because we only consider the API method calls. However, like all the detection schemes that are based on static analysis, Mal-Dozer is not resilient against dynamic code loading and reflection obfuscation, where the app downloads a malicious code and executes it at runtime. Moreover, MalDozer does not consider native codes. Finally, MalDozer uses many hyper-parameters for its deep learning model. The impact of using a lower number on accuracy detection and running time performance can be the subject of further investigations in future work. This work could also be extended to study the effect of a larger number of API method
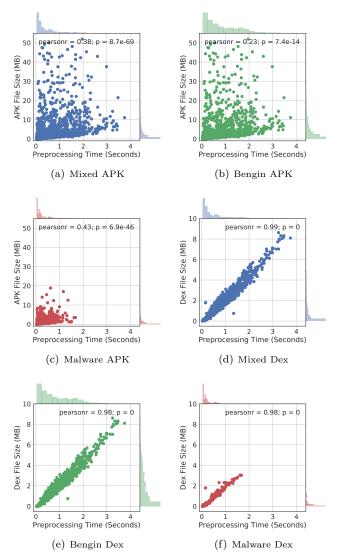
12

(a) Mixed APK

(b) Bengin APK

(c) Malware APK

(d) Mixed Dex

(e) Bengin Dex

(f) Malware Dex

Figure 10: Preprocessing Time vs APK and Dex Sizes (Server)

(a) Preprocess

(b) Prediction

Figure 11: Run-Time vs. Hardware

calls on detection accuracy, training time, and detection time. Furthermore, we plan to extend MalDozer to support the processing of native codes. In this context, Mal-Dozer could resist certain obfuscation techniques because we only consider the API call methods. However, Mal-Dozer is not immune to obfuscation (inherited from static analysis drawbacks) such as the use of reflection/dynamic code execution, as well as, other obfuscation techniques that rely on native code and non-framework service calls.

9. Related work

The Android malware analysis techniques can be classified as: *static analysis*, *dynamic analysis*, or *hybrid analysis*. The static analysis methods [24], [29], [30], [28], [31], [32], [33], [34], use static features that are extracted from the app, such as: requested permissions and APIs to detect malicious app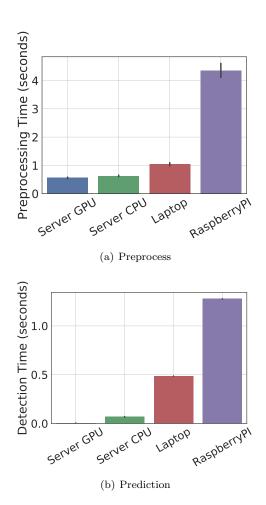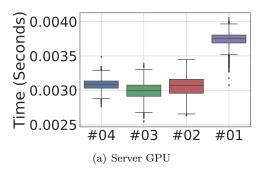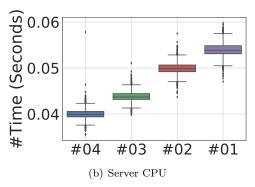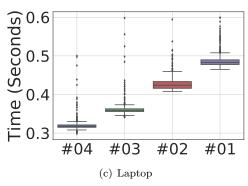. Some of these methods are generally not resistant to obfuscation. The dynamic analysis methods [35], [36], [37], [38],[39], [40], [41], [42] aim to identify behavioral signature or behavioral anomaly of the running app. These methods are more resistant to obfuscation. On the other hand, the dynamic methods offer limited scalability as they incur additional cost in terms of processing and memory. The hybrid analysis methods [43], [44], [45], [46], [47] combine between both analyses to improve detection accuracy, which costs additional computational cost. Assuming that malicious apps of the same family share similar features, some methods [48], [49], [50], [51], [52], [53], [54], [55], measure the similarity between the features of two samples (similar malicious code). Some methods [56], [57], [58] employ semantics-aware features such as control-flow graphs [59], data dependency graphs [60] and class dependence graphs [50]. The deep learning techniques are more suitable than conventional machine learning techniques for Android malware detection [43]. Research work on deep learning for Android malare detection are recently getting more attention [43], [61], [62], [63]. Differently from the existing deep learning solutions, MalDozer offers many advantages: (i) MalDozer provide automatic feature engineering for new types of malware
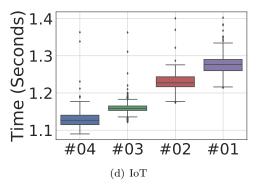
13

(a) Server GPU

(b) Server CPU

(c) Laptop

(d) IoT

Figure 12: Detection Time vs Model Complexity

in the training phase. (ii) MalDozer uses a minimal pre-processing, which fits small devices deployment. (iii) In addition to its high detection performance, MalDozer is able to attribute malware to its actual family with similar performance.

9.1. MalDozer vs. MAMADroid

In this section, we highlight the differences between MalDozer and MAMADroid [28], a recent and competitive malware detection solution that combines semantics-aware static features and machine learning. MAMADroid uses the sequence of abstracted API calls to build a behavioral model in the form of a Markov chain. The latter is used to extract features and perform classification. We compare our framework with MAMADroid with respect to the following points: (1) **Deployment:** Due to its large memory requirements, MAMADroid can only be deployed on high-power servers. In contrast, the detection component of MalDozer can efficiently run under multiple deployment architectures, including high-power and low-power machines as well as mobile and IoT devices. (2) **Design complexity:** The design of MAMADroid is complex since its preprocessing phase is composed of many steps. MAMADroid first constructs a call graph using flow-analysis tools such as: Soot [64] and FlowDroid [65]. Soot converts the DEX file into Jimple and FlowDroid is used for Taint analysis. Next, the constructed graph is used to generate the sequence of abstracted API calls, which are used to build a Markov chain. MalDozer, on the other hand, extracts the raw sequence of API method calls directly from the DEX file without any additional preprocessisng overhead. (3) **Adaptation to new APIs:** In MAMADroid, the release of new Android OS version with a new set of APIs implies that a manual intervention is needed to craft a new Markov model, and manually redefine the set of APIs to be fed to the classifier. MalDozer can automatically learn the new APIs directly from the DEX file, as they have a specific known format. (4)**Approach performance:** MAMADroid consumes a very large amount of memory (16GB of RAM) to extract the call graph from the app. Despite that, the graph extraction phase takes, in average, 25.4 seconds, and a maximum of 18 minutes. MalDozer is very lightweight. The average time to extract the features and classify the app in MalDozer is 5.3 seconds on the most resource-constrained IoT device. (5)**Ability to analyze apks:** The Soot tool, used by MAMADroid, fails to process 4.6% of the apps in the dataset of 43, 940 apps, which counts for more than 2, 000 apps. Such a failure does not happen in the case of MalDozer due to its simple preprocess. The above-discussed points clearly show the outperformance of MalDozer over MAMADroid.

10. Conclusion

We have presented MalDozer, an automatic, efficient and effective Android malware detection and attribution system. MalDozer relies on deep learning techniques and raw sequences of API method calls in order to identify Android malware. We have evaluated MalDozer on several small and large datasets, including *Malgenome*, *Drebin*, and our MalDozer dataset, in addition to a dataset of benign apps downloaded from Google Play. The evaluation

results show that MalDozer is highly accurate in terms of malware detection as well as their attribution to corresponding families. Moreover, MalDozer can efficiently run under multiple deployment architectures, ranging from servers to small IoT devices. This work represents a step towards practical, automatic and effective Android malware detection and family attribution.

References

[1] F. Delmastro, V. Arnaboldi, M. Conti, People-centric computing and communications in smart cities, IEEE Communications Magazine.

[2] A. Gilchrist, Industry 4.0: The Industrial Internet of Things, 2016.

[3] L. Yan, Y. Zhang, al, The Internet of Things: From RFID to the Next-Generation Pervasive Networked Systems, Auerbach Publications, 2008.

[4] Ericsson Mobility Report - https://tinyurl.com/gmnezg6 (2016).

[5] Smartphone os market share, 2017 q1 - https://tinyurl.com/y8tqgjfu (2017).

[6] Android Things on the Intel Edison board - https://tinyurl.com/gl9gglk (2016).

[7] Android Things OS - https://tinyurl.com/z3lf3ha (2016).

[8] RASPBERRY PI 3- https://tinyurl.com/ho7ngty (2017).

[9] Android Things - http://tinyurl.com/q5ko3zu (2016).

[10] Android Auto - http://tinyurl.com/hdsunht (2016).

[11] Android Wear - http://tinyurl.com/qfa55o4 (2016).

[12] RASPBERRY PI 2 - https://tinyurl.com/q65phuy (2017).

[13] MalGenome Dataset - http://tinyurl.com/combopx (2015).

[14] Drebin Dataset - http://tinyurl.com/pdsrtez (2015).

[15] Google Play - https://play.google.com/ (2016).

[16] T. Mikolov, I. Sutskever, al, Distributed Representations of Words and Phrases and their Compositionality, NIPS Neural Inf. Process. Syst.

[17] Android NDK - http://tinyurl.com/ppn559l (2016).

[18] Open Handset Alliance - https://tinyurl.com/2u76za (2016).

[19] Android Platform Architecture - https://tinyurl.com/hc7s4or (2017).

[20] Y. Kim, Convolutional neural networks for sentence classification, CoRR.

[21] J. Pennington, R. Socher, al., GloVe: Global Vectors for Word Representation, in: Conf. Empir. Methods Nat. Lang. Process., 2014.

[22] I. Goodfellow, Y. Bengio, Al, Deep Learning, MIT Press, 2016.

[23] Tensorflow - https://www.tensorflow.org (2017).

[24] D. Arp, M. Spreitzenbarth, H. Malte, H. Gascon, al., DREBIN: Effective and Explainable Detection of Android Malware in Your Pocket., in: Symp. Netw. Distrib. Syst. Secur., 2014.

[25] Contagiominidump - https://contagiominidump.blogspot.ca (2017).

[26] playdrone dataset - https://archive.org/details/playdrone-apks (2017).

[27] K. Allix, al, Androzoo: collecting millions of android apps for the research community, in: Proceedings of the 13th International Conference on Mining Software Repositories, MSR, 2016.

[28] E. Mariconti, L. Onwuzurike, P. Andriotis, E. De Cristofaro, G. Ross, G. Stringhini, Mamadroid: Detecting android malware by building markov chains of behavioral models, in: NDSS, 2017.

[29] Y. Feng, S. Anand, al., Apposcopy: Semantics-based detection of android malware through static analysis, in: Proceedings of the 22nd ACM International Symposium on Foundations of Software Engineering, 2014.

[30] W. Yang, J. Li, Y. Zhang, al., Apklancet: tumor payload diagnosis and purification for android applications, in: Proceedings of the 9th ACM symposium on Information, computer and communications security, 2014.

[31] Y. Zhongyang, Z. Xin, al., Droidalarm: an all-sided static analysis tool for android privilege-escalation malware, in: Proceedings of the 8th ACM SIGSAC symposium on Information, computer and communications security, 2013.

[32] E. B. Karbab, M. Debbabi, A. Derhab, D. Mouheb, Cypider: Building Community-Based Cyber-Defense Infrastructure for Android Malware Detection, in: ACM Computer Security Applications Conference (ACSAC), 2016.

[33] B. Sanz, I. Santos, al., Anomaly detection using string analysis for android malware detection, in: International Joint Conference SOCO13-CISIS13-ICEUTE13, Springer, 2014.

[34] E. B. Karbab, M. Debbabi, D. Mouheb, Fingerprinting Android packaging: Generating DNAs for malware detection, Digital Investigationdoi:10.1016/j.diin.2016.04.013.
URL http://linkinghub.elsevier.com/retrieve/pii/S1742287616300469

[35] G. Canfora, E. Medvet, al., Acquiring and Analyzing App Metrics for Effective Mobile Malware Detection, in: Proc. 2016 ACM Int. Work. Secur. Priv. Anal., 2016.

[36] E. M. B. Karbab, M. Debbabi, S. Alrabaee, D. Mouheb, DySign: Dynamic fingerprinting for the automatic detection of android malware, 2016 11th International Conference on Malicious and Unwanted Software, MALWARE 2016 (2016) 139–146arXiv:1702.05699, doi:10.1109/MALWARE.2016.7888739.

[37] M. Spreitzenbarth, F. Freiling, F. Echtler, al., Mobile-sandbox: having a deeper look into android applications, in: Proceedings of the 28th Annual ACM Symposium on Applied Computing, 2013.

[38] A. Ali-Gombe, al, Aspectdroid: Android app analysis system, in: ACM Conference on Data and Application Security and Privacy, 2016.

[39] Y. Zhang, M. Yang, B. Xu, Z. Yang, al., Vetting undesirable behaviors in android apps with permission use analysis, in: ACM conference on Computer and communications security (CCS), 2013.

[40] B. Amos, al, Applying machine learning classifiers to dynamic android malware detection at scale, in: Wirel. Commun. Mob. Comput. Conf., 2013.

[41] T.-E. E. Wei, C.-H. H. Mao, A. B. Jeng, H.-M. M. Lee, H.-T. T. Wang, D.-J. J. Wu, in: Trust. Secur. Priv. Comput. Commun., 2012.

[42] J. Huang, X. Zhang, al., Asdroid: Detecting stealthy behaviors in android applications by user interface and program behavior contradiction, in: Proceedings of the 36th International Conference on Software Engineering, 2014.

[43] Z. Yuan, Y. Lu, Z. Wang, Y. Xue, Droid-Sec: deep learning in android malware detection, in: ACM SIGCOMM Comput. Commun. Rev., 2014.

[44] M. Grace, Y. Zhou, al., Riskranker: scalable and accurate zero-day android malware detection, in: Proceedings of the 10th international conference on Mobile systems, applications, and services, 2012.

[45] S. Bhandari, R. Gupta, al., DRACO: DRoid analyst combo an android malware analysis framework, in: Proc. 8th Int. Conf. Secur. Inf. Networks, 2015.

[46] T. Vidas, J. Tan, J. Nahata, al., A5: Automated analysis of adversarial android applications, in: Proceedings of the 4th ACM Workshop on Security and Privacy in Smartphones & Mobile Devices, 2014.

[47] M. Lindorfer, M. Neugschwandtner, al., Andrubis–1,000,000 apps later: A view on current android malware behaviors, in: Building Analysis Datasets and Gathering Experience Returns for Security (BADGERS), IEEE, 2014.

[48] J. Kim, al, Structural information based malicious app similarity calculation and clustering, in: Proc. 2015 Conf. Res. Adapt. Converg. Syst., 2015.

[49] A. Ali-Gombe, al, Opseq: android malware fingerprinting, in: Proceedings of the 5th Program Protection and Reverse Engineering Workshop, 2015.

[50] L. Deshotels, V. Notani, A. Lakhotia, Droidlegacy: Automated familial classification of android malware, in: Proceedings of

ACM SIGPLAN on Program Protection and Reverse Engineering Workshop 2014, 2014.

[51] Y. Zhou, Z. Wang, al, Hey, you, get off of my market: detecting malicious apps in official and alternative android markets., in: NDSS, 2012.

[52] G. Suarez-Tangil, J. E. Tapiador, al., Dendroid: A text mining approach to analyzing and classifying code structures in android malware families, Expert Systems with Applications.

[53] H. Kang, J.-w. Jang, A. Mohaisen, H. K. Kim, Int. J. Distrib. Sens. Networks.

[54] Y.-D. Lin, Y.-C. Lai, C.-H. Chen, H.-C. Tsai, Identifying android malicious repackaged applications by thread-grained system call sequences, computers & security.

[55] P. Faruki, V. Laxmi, A. Bharmal, M. S. Gaur, V. Ganmoor, AndroSimilar: Robust signature for detecting variants of Android malware, J. Inf. Secur. Appl.

[56] M. Zhang, al., Semantics-aware android malware classification using weighted contextual api dependency graphs, in: Proceedings of the 2014 ACM SIGSAC Conference on Computer and Communications Security, 2014.

[57] M. Fan, J. Liu, X. Luo, K. Chen, al., Frequent subgraph based familial classification of android malware, in: Software Reliability Engineering (ISSRE), 2016 IEEE 27th International Symposium on, 2016.

[58] G. Meng, Y. Xue, Z. Xu, al, Semantic modelling of android malware for effective malware comprehension, detection, and classification, in: Proceedings of the 25th International Symposium on Software Testing and Analysis, 2016.

[59] M. Christodorescu, S. Jha, Seshia, Semantics-aware malware detection, in: IEEE Symposium on Security and Privacy (SP), 2005.

[60] M. Fredrikson, S. Jha, M. Christodorescu, R. Sailer, X. Yan, Synthesizing near-optimal malware specifications from suspicious behaviors, in: IEEE Symposium on Security and Privacy (SP), 2010.

[61] Z. Yuan, Y. Lu, Y. Xue, Droiddetector: android malware characterization and detection using deep learning, Tsinghua Science and Technology.

[62] S. Hou, A. Saas, al., Deep4maldroid: A deep learning framework for android malware detection based on linux kernel system call graphs, in: International Conference on Web Intelligence Workshops, 2016.

[63] S. Hou, A. Saas, al.s, Droiddelver: An android malware detection system using deep belief network based on api call blocks, in: International Conference on Web-Age Information Management, 2016.

[64] R. Vallée-Rai, P. Co, E. Gagnon, L. Hendren, P. Lam, V. Sundaresan, Soot-a java bytecode optimization framework, in: Proceedings of the 1999 conference of the Centre for Advanced Studies on Collaborative research, 1999.

[65] S. Arzt, S. Rasthofer, C. Fritz, al., Flowdroid: Precise context, flow, field, object-sensitive and lifecycle-aware taint analysis for android apps, Acm Sigplan Notices.

Appendix A. Confusion Matrices

2 Fold			3 Fold		
	PB[1]	PM		PB	PM
AB[2]	37,604	23	AB	36,882	745
AM	107	1,151	AM	40	1,218
5 Fold			10 Fold		
	PB	PM		PB	PM
AB	37,591	36	AB	37,612	15
AM	40	1,218	AM	44	1,214

[1] PB: Predicted Benign, PM: Predicted Malware
[2] AB: Actual Benign, PM: Actual Malware

Table A.17: Detailed Detection Result on Malgenome

2 Fold			3 Fold		
	PB[1]	PM		PB	PM
AB[2]	37,578	49	AB	37,435	192
AM	426	5,129	AM	233	5,322
5 Fold			10 Fold		
	PB	PM		PB	PM
AB	37,509	118	AB	37,457	170
AM	261	5,294	AM	168	5,387

[1] PB: Predicted Benign, PM: Predicted Malware
[2] AB: Actual Benign, PM: Actual Malware

Table A.18: Detailed Detection Result on Drebin

2 Fold			3 Fold		
	PB[1]	PM		PB	PM
AB[2]	37,246	381	AB	36,872	755
AM	1,421	18,668	AM	618	19,471
5 Fold			10 Fold		
	PB	PM		PB	PM
AB	36,779	848	AB	37,193	434
AM	436	19,653	AM	611	19,478

[1] PB: Predicted Benign, PM: Predicted Malware
[2] AB: Actual Benign, PM: Actual Malware

Table A.19: Detailed Detection Result on MalDozer

2 Fold			3 Fold		
	PB[1]	PM		PB	PM
AB[2]	36,673	954	AB	36,118	1,509
AM	1,821	31,245	AM	2,006	31,060
5 Fold			10 Fold		
	PB	PM		PB	PM
AB	36,621	1,006	AB	36,426	1,201
AM	1,585	31,481	AM	1,417	31,649

[1] PB: Predicted Benign, PM: Predicted Malware
[2] AB: Actual Benign, PM: Actual Malware

Table A.20: Detailed Detection Result on All Dataset

(a) Malgenome

(b) Drebin

(c) MalDozer